

# XenoControl

High Availability  
Management Console  
For Xen Clusters

Jérôme Petazzoni <[jp@enix.fr](mailto:jp@enix.fr)>

RMLL 2010

# Outline

- Who ? (please allow me to introduce myself)
- Why ? (goals & purposes)
- What ? (results & samples)
- How ? (tools & libraries)
- XenoControl internals
- Roadmap
- Questions

# Who

- Enix
  - Small hosting company
  - No money to buy expansive hardware
  - We love Open Source
  - Xen hosting since 2004
  - Before Xen : UML (UserModeLinux)

# Who

- SmartJog / TVRadio (TDF group)
  - Part of a big company
  - Have money, but don't want to waste it
  - They love Open Source
  - Streaming, transcoding, etc.
  - Resource usage varies all the time

# Why

- Need something to manage hosts & VM
  - List all VM in realtime
  - List resources usage and availability (CPU, RAM, storage, network)
  - Use standard commands & tools (xm, **LVM**, standard Xen config files)
  - Live migration of VM *without SAN*
  - Scriptability (GUI are out of question)



# Really, why ?

- Existing GUI also lack features
- Doing parallel SSH on 100+ dom0 is wrong
  - Needs some kind of registry/enumeration
  - Quickly turns into a hacky nightmare
- We want the system to be non intrusive
  - Must be able to plug/unplug XenoControl
  - Zero training to use the system (lazyness!)

# The real killer feature

- Live migration of VM with local storage
  - Why live migration ?
    - Redeploy resources
    - Hardware maintenance
  - Why local storage ?
    - Excellent performance without requiring high speed network (IB, 10G)
    - Cheap boxes available from misc. makers (4U, dual-socket Nehalem, 72GB RAM, 24x1 TB HDD, less than 10KEUR)

# What (Xen requirements)

- Standard Xen setup
- VM configuration stored in `/etc/xen/auto`
- Uses LVM on a single VG
  - VG naming must be the same everywhere
- Network setup is irrelevant
  - As long as you use standard Xen facilities (i.e. `/etc/xen/scripts/{vif,network}-*`)
  - All dom0 should be on the same Ethernet



# What (other requirements)

- Spread toolkit (more about this later)
- Python (2.6 and above)
- DRBD in the dom0 (not needed in domU)  
(but DRBD is not inherent to XenoControl)
- Linux dom0 !

# What (results)

```
milky:~# xenocontrol vmlist
```

```
Got a reply from #xenhost#xenlab
```

```
Got a reply from #xenhost#andromede
```

```
Got a reply from #xenhost#milky
```

```
Got a reply from #xenhost#medusa
```

host	name	vcpu	memMB	power_state
andromede	Domain-0	16	976	Running
andromede	enix.kran	1	1000	Running
andromede	zeen.obiwana	1	1000	Running
...	...	...	...	...
medusa	Domain-0	8	976	Running
medusa	enix.arachnee	1	512	Running
medusa	europnet.bisque	1	250	Running
...	...	...	...	...
milky	Domain-0	16	976	Running
milky	enix.dotcloud-1	1	4000	Running
milky	enix.dotcloud-2	1	1000	Running
milky	enix.dotcloud-3	1	1000	Running
milky	libe.back-dev	2	7000	Running

# What (more results)

milky:~# **xenocontrol hoststats**

```
...
  name    hvm    enabled  cpu    cpu_usage memMB  memfreeMB  nr_vm    disk_usage  net_usage
andromed True      True    16      5.6 % 73718  62975      8      0 / 209 KB   1 / 2 KB
medusa   True      True    8      130.1 % 32766  1242      38     36 / 2822 KB 177 / 1302 KB
milky    True      True    16      9.0 % 73719  11066     23     236 / 3030 KB 186 / 5919 KB
xenlab   True      True    4       2.1 % 8074   6841      2      91 / 14 KB   3 / 75 KB
```

4 host responded (4 enabled, 0 busy) for a total of 44 corethread hosting 71 vm

CPU : 42.5 free on 44 ; Used at 3.34 %

Memory : 80.2G free on 183.9G ; Used at 56.38 %

Disks : 7.7T free on 17.5T ; Used at 55.88 %

milky:~# **xenocontrol hostlist**

```
...
  name    hvm  cpu  memMB  freeMB  vm  stor_freeGB  xen_ver  kernel_version  cpu_model
andromed True  16  73718  62975  8      3986  3.2  2.6.30.1-xen-amd64 Intel(R) Xeon(R) CPU E5520 @ 2.27GHz
medusa   True  8  32766  1242  38      1625  3.2  2.6.30.1-xen-amd64 Intel(R) Xeon(R) CPU E5405 @ 2.00GHz
milky    True  16  73719  11066  23      2110  3.2  2.6.30.1-xen-amd64 Intel(R) Xeon(R) CPU E5520 @ 2.27GHz
xenlab   True  4  8074  6841  2      179  4.0  2.6.32-5-xen-amd64 Intel(R) Xeon(R) CPU X3430 @ 2.40GHz
```

4 host responded (4 enabled, 0 busy) for a total of 44 corethread hosting 71 vm

# What (easier management)

- All actions can use wildcards
- Need to do something on a bunch of VM ?  
# xenocontrol do\_something \*webfront\*  
# xenocontrol do\_something host17/\*sql\*
- Various hooks everywhere
  - Existing deployment tool was successfully integrated with Xenocontrol



# How (architecture decisions)

- Single-file Python script
  - Easy distribution even without packaging
  - Can upgrade code automatically
- Communication : spread toolkit
  - Reliable and efficient group communication
  - Setup is boring, but easy



# How (live migration steps)

- Prepare DRBD and remote LVM
- Magically switch block backend to DRBD
- Wait for DRBD to synchronize data
- Call “xm migrate”
- Magically switch block backend to LVM
- Tear down DRBD
- Tear down local LVM

# How (live migration steps)

Watch the lovely drawings  
(in separate file)

# Internals (automations)

- No centralized master
- Each complex command is an “automation” (recipe made of simple tasks)
- Automation = high-availability parallel job, run on the whole cloud of dom0 nodes
- Automation is controlled by *initiator*
- Automation can be resumed or aborted (in case of error or crash)

# Internals (requests)

- Automations can spawn workers (at least one, else no job is done)
- These workers will issue *requests*
- Requests can be executed by other hosts
- Requests are stateless
- State is kept on all nodes



# Internals (the magic)

How do we *magically* replace block devices ?

- “xm pause” the VM
- Use dmsetup to remap blocks (LVM trick)
- “xm unpause” the VM

The dmsetup call is quick :

- VM is not disturbed
- Unless running realtime tasks (VOIP...)



# Internals (restrictions)

- Works only for Linux dom0
  - But domU can be anything
  - Won't work with Solaris or NetBSD dom0 (But there might be another way!)
- Is not limited to DRBD and LVM
  - Can work with iSCSI, AOE ...
  - Should work with btrfs, glusterfs ...
  - Requirement : dmsetup

# Internals

## (black Python magic)

- Automations are made of workers
- Workers contain only logic control
- They are implemented using *continuations*
- Those *continuations* are implemented using generators and special *yield* syntax
- All real job is done by issuing *requests*
- *Requests* can execute elsewhere

# Internals

## (internals of internals)

- Each automaton = 1 spread group
- Each group = one master (initiator) elected
- Only the master will send data
- Other nodes will run the control logic, receive all data, but send nothing
- State is kept in sync on all nodes
- Allow for recovery in case of failure

# Internals

## (black magic unveiled, part 1)

- Engine code :  
`start_at_least_one_worker()`
- Worker code :  
... some control logic  
... then need to do something  
`response = yield Request(requestparams)`  
... some more control logic  
(hopefully using the response)



# Internals

## (black magic unveiled, part 2)

- Engine (pseudo-)code :

```
# For each worker...
request = worker.next() # first request
while not_finished:
    # initiator is the only one to send
    if initiator:
        send_to_network(request)
    # but everyone receives requests
    # and their responses, to keep state
    response = read_from_network()
    request = worker.send(response)
# And multiple workers run in parallel :)
```



# Internals

## (black magic by the book)

```
class DeployAutomation(AutomationProcessor):
    def s_init(self, args):
        vm_to_deploy = parse_opts(args)
        for vm in vm_to_deploy:
            possible_hosts = Request('can_host', vm, host='*')
            chosen_host = self.random.choice(possible_hosts)
            Request('set_busy', True, host=chosen_host)
            self.worker.add(self.s_deploy(vm, chosen_host))
        def s_deploy(self, vm, chosen_host):
            result = Request('install_vm', vm, host=chosen_host)
            if result != 'OK':
                Request('cleanup_vm', vm, host=chosen_host)
            else:
                Request('start_vm', vm, host=chosen_host)
            Request('set_busy', False, host=chosen_host)
```

# Internals (limitations)

- Can't use random
  - Initialize each worker with a common seed
  - So you can use Python's random after all
- Can't interact with outer world (sockets...)
  - Workers must delegate communication
  - State must be consistent :  
inform other nodes of what's happening  
... this is actually done automatically

# Roadmap (what we want to do)

- Code cleanup
- Allow VG with different names
- Allow multiple VG
- Implement migration from/to iSCSI
- Add external control web server

# Roadmap (what you can do)

- If you know kung-fu : new automations
  - Support for AOE, iSCSI, glusterfs, btrfs ...
  - Integration with SAN management
- If you know Python : better output format
  - Add proper tabular/XML/JSON output
- If you know english : better messages
  - Right now everything is in frenglish



# Roadmap

## (cool stuff we dream about)

- Integration with ovirt
  - So you can get all the cool ovirt tools
- KVM / OpenVZ support
  - Because Xen sometimes sucks
  - And because KVM and OpenVZ are great
- Replace spread with some AMQP flavour
  - Spread has rough edges, but gracefully handles every nuke we throw at it



# Conclusion

- Useful tool, used internally by two really different user profiles
- Neat architecture  
(or at least, that's what we think)
- Needs a great amount of cleanup, doc ...
- You can also see XenoControl as a POC

# Questions ?

<http://xenocontrol.enix.org/>

<http://www.enix.fr/>  
[jp@enix.fr](mailto:jp@enix.fr)

<http://www.smartjog.com/>